

The automated testing of randomness with multiple statistical batteries

Ľubomír Obrátil
lubomir.obratil@gmail.com

22. 6. 2017

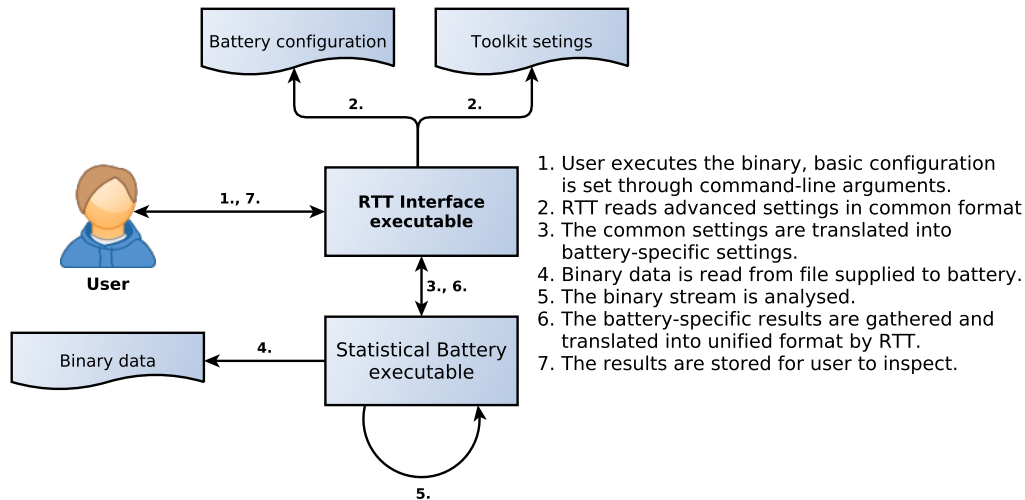
Thesis structure

- Creation of a unified interface supporting multiple statistical batteries.
 - NIST Statistical Testing Suite
 - Dieharder
 - TestU01
- Conducting the baseline (control) experiment to create a reference point for the further experiments.
- Evaluating randomness of outputs of well-known cryptographic primitives.
- Analysing validity of the Dieharder battery results.

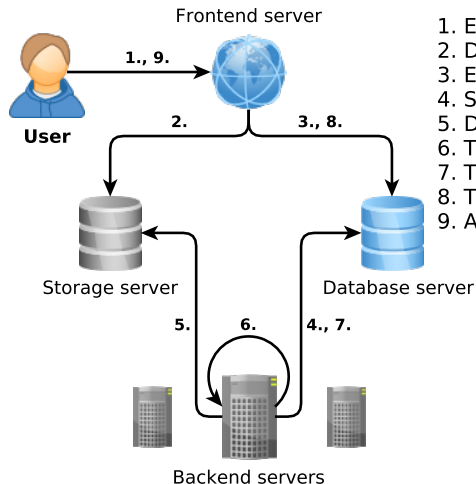
Randomness Testing Toolkit – overview

- Design and implementation of a tool for consistent randomness evaluation.
- Intended to be used by users without prior knowledge about statistical testing as well as by researchers in CROCS.
- The developed tool (RTT) acts as a interface between the user and the statistical batteries – common format of the battery settings and results.
- The toolkit supports multiple statistical batteries – NIST STS, Dieharder and TestU01; it is possible to add more batteries over time.
- Both standalone program and online service were developed.

Randomness Testing Toolkit – local interface



Randomness Testing Toolkit – web service



Statistical testing of randomness – 1/2

Testing hypothesis – H_0

During the experiments, we evaluated the hypothesis that the analysed data were produced by a truly random generator. We denote the hypothesis as H_0 (null hypothesis).

Statistical battery

Software with the purpose of detecting biases in data stream; collection of statistical tests.

Statistical test

Single unit in a statistical battery, checks some property of the data; e.g. count of ones. Output of a test is the probability that the analysed data were produced by TRNG. Each test in a given battery will either fail (H_0 rejection) or pass (H_0 retainment).

Statistical testing of randomness – 2/2

Significance level – α

The significance level is set prior to the experiments (usually 0.001) and based on it, the null hypothesis is rejected or retained.

False positive (Type I error)

The false positive result is observed when H_0 holds true but it is rejected – stream produced by TRNG is evaluated as non random. Probability of Type I error is α .

False negative (Type II error)

The false negative result is observed when H_0 is false but it is not rejected – stream generated by biased generator is evaluated as random.

Establishing baseline results

- The result of a battery is obtained from the proportion of the failed tests in the battery (e.g. we consider the data biased if more than 2 out of 15 tests fail). However, even data generated by a perfect TRNG may fail some tests (false positives).
- To identify the maximum count of failed tests (assuming that the H_0 holds) we processed 8 TB of data generated by a quantum random generator (control experiment). Further experiments were interpreted based on these results.
- In order to make the result interpretation more straight-forward, results of certain tests were grouped together. As a consequence, the counts of failed tests are closer to expected (theoretical) numbers and therefore more easy to evaluate.

Analysis of well-known algorithms

- Analyse outputs of well-known cryptographic algorithms (AES, DES, RC4, etc.).
- Observe the security margins of the algorithms.
- Compare the results to other approach in CRoCS.

Analysed algorithms

- In total, 72 different data streams were analysed.
- The data streams were outputs from 16 distinct round-reduced cryptographic algorithms.
- The algorithms were chosen based on their popularity (AES, DES, RC4, ...) or their success in crypto competitions eSTREAM and SHA3 (Rabbit, Keccak, Grøstl, ...).

Analysis conditions

- Each datastream was 8GB long.
- The conditions of analysis were same as in the previous experiment.
- The interpretation of the result was based on the results of the baseline

Analysis of the results of Dieharder battery

- Examine behavior of Dieharder battery during truly random data analysis.
- Analyse the distribution of the results.

Usable testbed analysis – 2/3

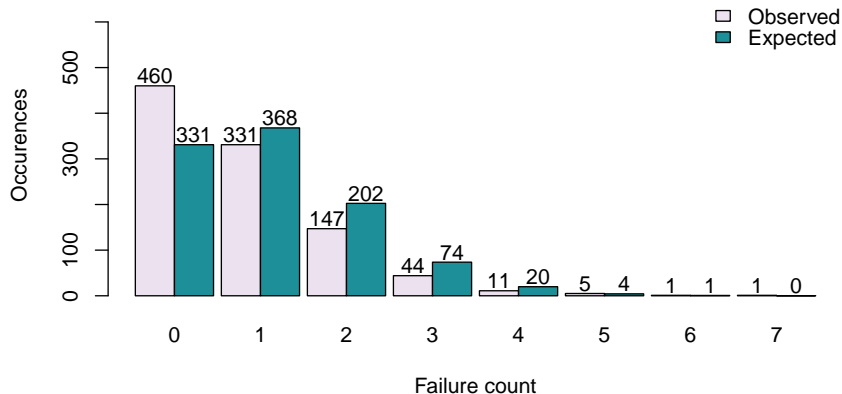
Algorithm	Biased round RTT	Biased round EACirc	Security Margin
AES	3	3	7 – 70%
BLAKE	1	1	15 – 93.7%
Grain	6*	2	7 – 53.8%
Grøstl	2	2	12 – 85.7%
HC-128	–	–	0 – 100%
JH	6	6	36 – 85.7%
Keccak	3	2	21 – 87.5%
MD6	10*	8	94 – 90.3%
Rabbit	4*	0	0 – 0%
RC4	0*	–	0 – 0%
Salsa20	2	2	18 – 90%
SINGLE-DES	5	4	11 – 68.7%
Skein	4	3	68 – 94.4%
SOSEMANUK	4	4	21 – 84%
TEA	5	4	27 – 84.3%
TRIPLE-DES	3	2	13 – 81.2%

References

- **Randomness Testing Toolkit**
<https://github.com/crocs-muni/randomness-testing-toolkit>
- **EACirc**
<https://github.com/crocs-muni/eacirc>
- **NIST Statistical testing suite**
http://csrc.nist.gov/groups/ST/toolkit/rng/documentation_software.html
- **Dieharder**
<http://www.phy.duke.edu/~rgb/General/dieharder.php>
- **TestU01**
<http://simul.iro.umontreal.ca/testu01/tu01.html>

Baseline experiment – uncorrected results

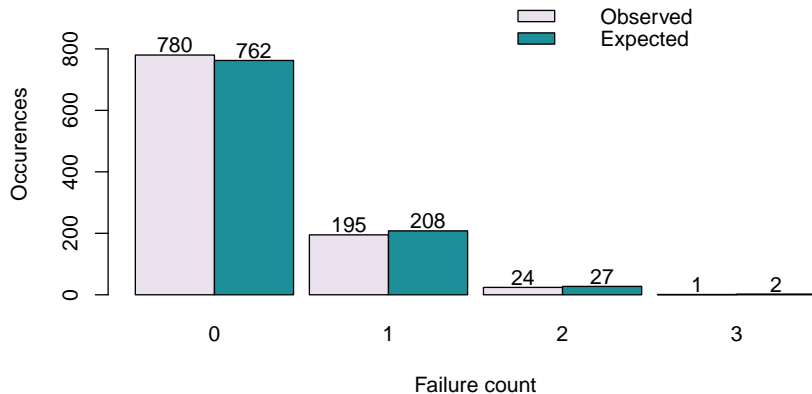
Dieharder (original), 110 tests
 χ^2 statistic p-value = 5.32e-17



Baseline experiment – corrected results

Dieharder (corrected), 27 tests

χ^2 statistic p-value = 0.382



Baseline experiment – resulting reference

Battery name	Closeness to expected results		Allowed failures
	Uncorrected	Uncorrected	
Dieharder	$5.32 \cdot 10^{-17}$	0.38	3/27
NIST STS	$2.17 \cdot 10^{-2}$	$4.44 \cdot 10^{-7}$	2/15
TU01 Small Crush	0.71	0.95	2/10
TU01 Crush	$3.36 \cdot 10^{-11}$	$3.31 \cdot 10^{-3}$	3/32
TU01 Rabbit	$2.02 \cdot 10^{-5}$	$1.45 \cdot 10^{-23}$	2/16
TU01 Alphabit	$2.14 \cdot 10^{-8}$	$2.8 \cdot 10^{-7}$	1/4
TU01 Block Alphabit	$1.87 \cdot 10^{-68}$	$5.15 \cdot 10^{-47}$	1/4

Notable results

- **Grain** – Tests smarsa_MatrixRank and scomp_LinearComp (Crush, Rabbit) will fail in 3, 4, 5 and 6-round configuration.
- **MD6** – Tests smarsa_MatrixRank and sspectral_Fourier3 (Crush, Rabbit) will fail in 9 and 10-round configuration.
- **Rabbit** – Tests sstring_HammingIndep and sstring_PeriodsInStrings (Crush, Rabbit, Alphabit, Block Alphabit) will fail in **full** configuration.
- **RC4** – Tests sknuth_SimpPoker and sknuth_Gap (Crush) will fail in **full** configuration.

Analysis of Dieharder – 1/2

Analysed data

- 8TB of quantum random data processed continuously by the tests - single application of a test to a data stream will yield single first-level p-value
- Uniformity of the first level p-values was analysed.
- The p-values should be uniformly distributed on the interval $< 0, 1 >$.
- Total of 110 sets of p-values (single set per raw, uncorrected test) was inspected.
- Each set had a different size – usually between 1 to 2 millions of p-values per set.

Experiment results

- Out of 110 p-value sets, 39 sets were not uniformly distributed
- Chi-Square (χ^2) statistic used for uniformity testing. When the p-value of the statistic was less than 0.001, the inspected set was considered non-uniform.
- Flawed non-uniform distributions can have impact on Dieharder results.

Analysis of Dieharder – 2/2

